

DATA304 Big Data Analysis


6. Text Classification 2: Learning with Limited Labels 1

SeongKu Kang

Korea University

<https://www.idea.korea.ac.kr/>


We are opening an e-commerce platform!



We have a wide variety of items, with new arrivals every day!



Nike Air Force is a range of athletic shoes made by Nike. It was created by designer Bruce Kilgore and was the first basketball shoe ...




Harry Potter and the Philosopher's Stone is a fantasy novel written by British author J. K. Rowling. Harry Potter lives with his abusive uncle and aunt, Vernon ...



Lover is the seventh studio album by the American singer-songwriter Taylor Swift. Singles from Lover: "Me!", ... , "Lover", "The Man", "Cruel Summer" ...


We are opening an e-commerce platform!




Your platform should be able to categorize a wide range of items (**classification**).



Clothing, Shoes & Jewelry
> Men > Shoes
> Fashion Sneakers



Books > Teen & Young Adult
> Literature & Fiction
> Action & Adventure
> Fantasy



CDs & Vinyl
> Pop > Vocal Pop


Text classification

- Why Important?

- Core task to analyze your data: assign predefined labels (e.g., brand, sentiment, topic).
 - Widely used in search engines, recommender systems, spam detection, and more.
 - Forms the basis for many advanced applications such as QA, dialogue systems, and personalization.



Clothing, Shoes & Jewelry
> Men > Shoes
> Fashion Sneakers



Books > Teen & Young Adult
> Literature & Fiction
> Action & Adventure
> Fantasy



CDs & Vinyl
> Pop > Vocal Pop


- Our learning path:

- We have represented texts as vectors.
- Next, we will conduct **classification** using these representations.
- We will mainly follow the **pretrain-and-fine-tune paradigm** to adapt *general language knowledge* to *specific classification tasks*.

Recap: Pretraining + Fine-tuning

- **Pretraining**


- **Motivation:** there are abundant text data on the web, with rich information of linguistic patterns and knowledge about the world
- **Goal:** Learn **general-purpose representations** without human labeling
- **Example:**
 - Word2Vec (predict neighboring words)
 - BERT (masked language modeling, next sentence prediction)



Recap: Pretraining + Fine-tuning

- **Fine-tuning**

- **Motivation:** pretrained models capture general knowledge, but tasks like classification require task-specific knowledge
- **Goal:** **adjust the pretrained model parameters** to perform well on a target (downstream) task
 - Requires *much less labeled data* than training from scratch
 - Leverages *knowledge learned during pretraining*
- **Examples:**
 - Sentiment classification using BERT




- Downstream task**
- Spam mail classification
 - Sentiment classification
 - Search
 - Question and answering
 - ...

Recap: Pretraining + Fine-tuning

- Downstream task: sentiment classification
 - Predict whether a given sentence is positive, neutral, or negative

- How it works:
 - Given a **pretrained** BERT, we represent each sentence as a vector
 - Popular choice: CLS representation, average of all contextual embeddings
 - We add a small classification head (a small network)
 - The model is then **fine-tuned** with labeled data

- Training choice:
 - A. Full fine-tuning: update all parameters
 - B. Partial fine-tuning: update a subset of parameters
 - Freeze the encoder + update only the head
 - Freeze most of the encoder + update the top layers and the head



Recap: Classification task

- **Classification type & loss:**

- Binary classification
- Multi-class classification
- Multi-label classification


$$\mathcal{L}_{\text{BCE}}(\theta) = - \sum_{n=1}^N \left[y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n) \right]$$

$$\mathcal{L}_{\text{CE}}(\theta) = - \sum_{n=1}^N \sum_{c=1}^C \mathbf{1}[y_n = c] \log \hat{y}_{n,c}$$

$$\mathcal{L}_{\text{Multi-label}}(\theta) = - \sum_{n=1}^N \sum_{c=1}^C \left[y_{n,c} \log \hat{y}_{n,c} + (1 - y_{n,c}) \log(1 - \hat{y}_{n,c}) \right]$$

- **Think about the whole picture!**

- Each input text is represented as a vector using a pretrained model.
- By minimizing the loss, **you fine-tune the model** (either full vs. partial depending on your choice)



Recap: Classification evaluation

- **Metrics:**


- Accuracy: Overall fraction of correct predictions.
- Precision: Among predicted positives, how many are correct.
- Recall: Among actual positives, how many are found.
- F1-score: Harmonic mean of Precision and Recall.

	urgent	normal	spam	
urgent	8	10	1	$\text{precision}_u = \frac{8}{8+10+1}$
normal	5	60	50	$\text{precision}_n = \frac{60}{5+60+50}$
spam	3	30	200	$\text{precision}_s = \frac{200}{3+30+200}$
	$\text{recall}_u = \frac{8}{8+5+3}$	$\text{recall}_n = \frac{60}{10+60+30}$	$\text{recall}_s = \frac{200}{1+50+200}$	

- **Averaging methods:**

- Macro-averaging: Compute metrics per class, then average.
 - Treats all classes equally, regardless of size.
- Micro-averaging: Aggregate TP, FP, FN across classes, then compute metrics.
 - Weighted by class frequency, favors larger classes.
- Use Macro when class imbalance is a concern; Micro when overall performance is more important.

Course plan



Sparse representation

Dense static representation

Dense contextual representation

Representing texts as vectors

Classification task

Learning with limited labels

Text classification

Lexical retrieval

Dense retrieval

LLM-enhanced IR

Search

Node representation


Graph neural networks

Leveraging relations

Text mining & Information retrieval


Learning with limited labels


- So far, we have discussed how to train classifiers using labeled data.
 - For every input text x , a corresponding class label y is provided.
- *But in reality, do we always have labels for all data instances?*
- We will explore how to deal with the scarcity of labeled data.



Learning with limited labels

- So far, we have discussed how to train classifiers using labeled data.
 - For every input text x , a corresponding class label y is provided.
- *But in reality, do we always have labels for all data instances?*
- We will explore how to deal with the scarcity of labeled data.
 1. **Semi-supervised learning**
 - “*How can we effectively leverage unlabeled data?*”
 2. **Multi-task learning**
 - “*If one task doesn't have enough labels, can we borrow signals from related tasks?*”
 3. **Adversarial learning**
 - “*What if labeled and unlabeled data come from different distributions?*”





(Figure credits): many figures in this section are borrowed from Mr. Baixu Chen's slides and Dr. Kevin Clark's slides

Semi-supervised learning

Learning types according to label availability

- **Supervised learning**

$$\mathcal{D}_l = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$

Sentiment labels	
Positive	Positive
Negative	Negative
Neutral	Neutral
:	:
	Negative

- All training samples have **ground-truth labels**
- **Limitation:** labeling requires **human effort** (often expert knowledge), making it infeasible to obtain very large labeled datasets.

- **Semi-supervised learning**

$$\mathcal{D}_l = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$


$$\mathcal{D}_u = \{x_1, x_2, \dots, x_m\}$$

Sentiment labels	
Positive	Positive
Negative	Negative
?	?
:	:
	?

- Only a small set of samples have ground-truth labels, i.e., $m \gg n$
- **Advantage:** Collecting unlabeled samples is *much easier* compared to human-labeled data


Semi-Supervised Learning (SSL)

- **Goal:** Use both labeled and unlabeled data during training
 - Labeled data provides **direct supervision**
 - Unlabeled data helps the model **generalize better** when used with the right techniques



Semi-Supervised Learning (SSL)

- **Goal:** Use both labeled and unlabeled data during training
 - Labeled data provides **direct supervision**
 - Unlabeled data helps the model **generalize better** when used with the right techniques



Semi-Supervised Learning (SSL)

- **Goal:** Use both labeled and unlabeled data during training
 - Labeled data provides **direct supervision**
 - Unlabeled data helps the model **generalize better** when used with the right techniques


- **Key assumptions:**

1. Smoothness assumption

- Nearby data points in the input space should have the same label
- “*If they look similar, they are similar*”

2. Low-density assumption

- A good decision boundary between classes should avoid regions with many data points.
- “*Boundaries should pass through low-density areas*”



Semi-Supervised Learning (SSL)


- Two main (and foundational) approaches we will cover:

1. Pseudo labeling

- Let the model *assign labels to unlabeled data* using its *confident predictions*.
- Confident predictions usually match what the model has already learned from similar labeled examples

2. Consistency regularization



- For *the same input with small changes*, the model should make *consistent prediction*.
 - E.g., Adding small noise to vectors, or a minor word change.
- Encourages decision boundaries to lie in low-density regions.



The model should make the same prediction for any point in the circle

1. Pseudo labeling

- **Key idea:** let the model *assign labels to unlabeled data* using its *confident predictions*.
- During training with labeled data, the model becomes confident about certain unlabeled samples.
 - This confidence comes from having already seen similar labeled examples.
 - Such predictions can be reused as “pseudo-labels” for training.




Between the two softmax outputs,
which one would you trust more, and why?

1. Pseudo labeling: process


(Figure credit): Deep learning bible

- **Step 1: Train with labeled data**
 - Start with a small portion of labeled data
 - Train the model in a supervised way
- **Step 2: Predict labels for unlabeled data**
 - Apply the trained model to unlabeled data
 - Obtain predictions with varying confidence levels (max prob.)
- **Step 3: Generate pseudo-labeled data**
 - Keep only high-confidence predictions as pseudo-labels
 - Treat them as if they were true labeled samples
- **Step 4: Retrain the model**



1. Pseudo labeling: self-training

- By learning with pseudo-labels, the model gradually gains more knowledge about unlabeled data.
- **Self-training:** Repeat the pseudo-labeling process to propagate this knowledge and expand supervision.
 - Self-training == Iterative pseudo-labeling




1. Pseudo labeling: self-training

- By learning with pseudo-labels, the model gradually gains more knowledge about unlabeled data.
- **Self-training:** Repeat the pseudo-labeling process to [propagate this knowledge](#) and expand supervision.
 - Self-training == Iterative pseudo-labeling

Algorithm 1 Self-training with Pseudo-Labeling

- 1: **Input:** Labeled dataset \mathcal{D}_l , Unlabeled dataset \mathcal{D}_u , Confidence threshold τ
- 2: **Output:** Trained model f
- 3: Train base model f on \mathcal{D}_l
- 4: **while** not converged **do**
- 5: Predict pseudo-labels for samples in \mathcal{D}_u using f
- 6: Select high-confidence predictions: $\mathcal{D}_p = \{(x, \hat{y}) \mid \max f(x) \geq \tau\}$
- 7: Update training set: $\mathcal{D}_l \leftarrow \mathcal{D}_l \cup \mathcal{D}_p$
- 8: **Update** f on the expanded \mathcal{D}_l
- 9: **end while**
- 10: **return** f



1. Pseudo labeling: summary

(Figure credit) IceBerg: Debiased Self-Training for Class-Imbalanced Node Classification, WWW'25

- Pseudo labeling is a simple and effective way to leverage unlabeled data.
- However, model predictions are not as accurate as human annotations.
 - In the worst case, pseudo-labeling can even perform worse than using only labeled data.



- With limited labels for minority class (orange), the model struggles to make accurate predictions across the whole space.
- Pseudo labels may only capture partial patterns of the unlabeled data.

- Limitations:
 1. **Error propagation:** Wrong pseudo-labels reinforce mistakes in training
 2. **Overconfidence:** The model can become overly confident, even in its wrong predictions

Semi-Supervised Learning (SSL)

- Two main (and foundational) approaches we will cover:


1. Pseudo labeling

- Let the model *assign labels to unlabeled data* using its *confident predictions*.
- Confident predictions usually match what the model has already learned from similar labeled examples

2. Consistency regularization



- For *the same input with small changes*, the model should make *consistent prediction*.
 - E.g., Adding small noise to vectors, or a minor word change.
- Encourages decision boundaries to lie in low-density regions.




The model should make the same prediction for any point in the circle

2. Consistency regularization

- Key idea: for *the same input with small changes*, the model should make *consistent prediction*.
 - Add small perturbations (e.g., noise, augmentation) → prediction should remain stable
 - Equivalently, nearby data points should get similar predictions

● ● : labeled data
● : unlabeled data



2. Consistency regularization


- Key idea: for *the same input with small changes*, the model should make *consistent prediction*.
 - Add small perturbations (e.g., noise, augmentation) → prediction should remain stable
- Instantiation:

$$\mathcal{L} = \mathcal{L}_{sup}(D_l) + \lambda \mathcal{L}_{cons}(D_u)$$

$$\mathcal{L}_{cons} = \mathbb{E}_{x \in D_l \cup D_u} \left[\|f(x; \theta) - f(\tilde{x}; \theta)\|^2 \right]$$

x: original input, \tilde{x} : perturbed input

- How to add perturbations?
 - Add small random noise to the embeddings
 - Word dropout or masking
 - Data augmentation (e.g., rotation, crop in vision)
 - ... and many other ways



2. Consistency regularization: insights

✓ Regularization prevents overfitting by discouraging models from memorizing data


- Instantiation:

$$\mathcal{L} = \mathcal{L}_{sup}(D_l) + \lambda \mathcal{L}_{cons}(D_u)$$

$$\mathcal{L}_{cons} = \mathbb{E}_{x \in D_l \cup D_u} \left[\|f(x; \theta) - f(\tilde{x}; \theta)\|^2 \right]$$

x : original input, \tilde{x} : perturbed input

- Insight 1: it forces the *predicted probability distributions* to be similar for original and perturbed inputs.



- The model tends to be naturally less confident on perturbed input
 - As a result, the model has lower confidence overall
- ✓ **Regularization effect:** prevents the model from becoming overconfident, which happens when it memorizes the data (overfitting)

2. Consistency regularization: insights


- Instantiation:

$$\mathcal{L} = \mathcal{L}_{sup}(D_l) + \lambda \mathcal{L}_{cons}(D_u)$$

$$\mathcal{L}_{cons} = \mathbb{E}_{x \in D_l \cup D_u} \left[\|f(x; \theta) - f(\tilde{x}; \theta)\|^2 \right]$$


x: original input, \tilde{x} : perturbed input

- **Insight 2:** unlike pseudo labeling that relies only on *high-confidence samples*, it uses *all unlabeled data*.
 - This allows leveraging much larger portions of the dataset.



+, ▽: labeled data

Pseudo labeling



Consistency regularization


2. Consistency regularization: temporal ensemble

- One practical issue: during training, the model often makes **unstable prediction for unlabeled data**.
 - Predictions fluctuate across epochs, even without perturbations
 - With perturbations, instability becomes even worse → This makes whole training unstable!

$$\mathcal{L}_{cons} = \mathbb{E}_{x \in D_l \cup D_u} \left[\|f(x; \theta) - f(\tilde{x}; \theta)\|^2 \right]$$

For unlabeled data, we don't have gold labels.
Model prediction is often highly unstable!

[Softmax output from the model]




2. Consistency regularization: temporal ensemble

- How to stabilize the model prediction?

✓ Model ensemble

- Instead of using a single model, train multiple models and aggregate their predictions
- **Why?**
 - Different initializations → different converged parameters.
 - By averaging predictions, **variance is reduced**, making **more stable & reliable prediction!**




2. Consistency regularization: temporal ensemble

- How to stabilize the model prediction?

✓ Temporal ensemble

- Independently training multiple models is expensive. *Can we get the similar effect with just one model?*
- **Key idea:** Aggregate predictions of the same model over different training epochs
- Averaging over time reduces fluctuations. The prediction becomes more **stable & reliable!**



2. Consistency regularization: temporal ensemble


✓ Temporal ensemble

- **Key idea:** Aggregate predictions of the same data over different training epochs
- **Instantiation:** Conduct ensemble by **exponential moving average (EMA)** of model parameters


$$\theta_{te}^{(t)} \leftarrow \alpha \theta_{te}^{(t-1)} + (1 - \alpha) \theta^{(t)}$$

$\theta^{(t)}$: model parameter at epoch t
 $\theta_{te}^{(t)}$: temporal ensemble at epoch t

α : a hyperparameter that controls
the degree of past knowledge



- Temporal ensemble aggregates past model parameters during training.
- It *follows the latest model slowly*, while remaining much more *stable by averaging predictions across epochs*.



Temporal ensemble update process

2. Consistency regularization: temporal ensemble

- **Consistency regularization with temporal ensemble**
 - We use $f(x; \theta_{te})$, the prediction from temporal ensemble
 - Provides a **more reliable and stable guidance** for unlabeled data

$$\mathcal{L} = \mathcal{L}_{sup}(D_l) + \lambda \mathcal{L}_{cons}(D_u)$$


$$\mathcal{L}_{cons} = \mathbb{E}_{x \in D_l \cup D_u} \left[\|f(x; \theta_{te}) - f(\tilde{x}; \theta)\|^2 \right]$$

Algorithm 2 Consistency Regularization with Temporal Ensemble

- 1: **Input:** Labeled dataset D_l , Unlabeled dataset D_u
 - 2: **Output:** Trained model $f(\cdot; \theta)$
 - 3: Initialize parameters θ , ensemble parameters $\theta_{te} \leftarrow \theta$
 - 4: **for** epoch $t = 1$ to T **do**
 - 5: Compute supervised loss on D_l : \mathcal{L}_{sup}
 - 6: Compute consistency loss on D_u : $\mathcal{L}_{cons} = \|f(x_u; \theta_{te}) - f(\tilde{x}_u; \theta)\|^2$
 - 7: Update θ by minimizing $\mathcal{L} = \mathcal{L}_{sup} + \lambda \mathcal{L}_{cons}$
 - 8: Update ensemble parameters: $\theta_{te}^{(t)} \leftarrow \alpha \theta_{te}^{(t-1)} + (1 - \alpha) \theta^{(t)}$
 - 9: **end for**
 - 10: **Return:** $f(\cdot; \theta)$
-


The original version

$$\mathcal{L}_{cons} = \mathbb{E}_{x \in D_l \cup D_u} \left[\|f(x; \theta) - f(\tilde{x}; \theta)\|^2 \right]$$



2. Consistency regularization: summary

- For the same input with small changes, the model should make consistent predictions!
- Pros:
 - **Effective use of unlabeled data:** requires only perturbations and applied to broad unlabeled data
 - **Simple concept:** easy to add as a regularizer to existing models
- Cons:
 - **Perturbation sensitivity:** effectiveness depends on the choice of noise/augmentation.
 - Researchers still actively study this topic.
 - **Weak direct guidance:** does not provide explicit labels, so boundaries may still be uncertain.




Semi-Supervised Learning (SSL): summary

- Two main (and foundational) approaches we covered:
 1. **Pseudo labeling**
 - **Strength:** Provides a direct training signal (like ground truth)
 - **Limitation:** Can amplify errors — if the model is overconfident in wrong predictions.
 2. **Consistency regularization**
 - **Strength:** Acts as a regularizer that prevents the model from becoming overconfident.
 - **Limitation:** Provides weak guidance since no explicit labels are given.
- They have complementary relationship!
 - *Pseudo-labeling provides strong but noisy signals, while consistency regularization stabilizes them.*
 - *Almost all state-of-the-art SSL combines two approaches as a holistic framework!*

(Optional) Overview of SSL

- In this class, we covered up to Mean Teacher, published in 2017.
- Almost all subsequent work are based on two approaches we studied.



Recommended readings

- **Articles:**
 - The Universe of AI – Semi-supervised learning
 - Korean: <https://wikidocs.net/255169>
 - English: <https://wikidocs.net/255197>
- **Papers:**
 - Consistency regularization with temporal ensemble
 - Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised deep learning results, NeurIPS'17
 - (Purely optional, beyond the scope of this class)
 - Debiased Self-Training for Semi-Supervised Learning, NeurIPS'22